

Assumptions

1. There are **P** processors in CPU architecture
2. The order in which edges are entering the incidence matrix does not matter
3. Storing of the matrix in memory is in row-major format

Observations

1. First, we will take the transpose of the Incidence Matrix so that we can traverse row wise to optimize cache performance
 - a. Each processor will get M/P rows where $M = \text{\#Edges}$
2. Each row will now update one cell of `adj_mat` which will be shared between all the processors
 - a. If there is an overlap while updating the cell critical section will handle it

Pseudo Code for the Parallel Algorithm

```
// Global Region
M ← #Edges in graph
P ← #Processors
Local_size ←  $M/P$ 
Adj_mat ← Zeroes(N,N)
Inc_mat ← transpose(Inc_mat)

//Adj_mat and Inc_mat are shared between all the processors

// Parallel Region
// Parallelise this loop into processors with  $M/P$  rows to each one
LOOP (i,0,M) :
    X = -1
    Y = -1
    LOOP (j,0,N) :
        if(Inc_mat[i][j]>0)
            if(X<0)
                X ← j
            else
                Y ← j
    // Critical Section
    if(Y<0)
        Adj_mat[X][X] ← Adj_mat[X][X] + 2
    else
        Adj_mat[X][Y] ← Adj_mat[X][Y] + 1
```

Drawbacks

1. While updating one cell of `Adj_mat` there can be overlaps such that two processors are accessing the cell simultaneously, in this case, the performance of algorithm will be affected